

# Kubernetes安装部署操作手册

作者姓名：杨遥

项目组：中科软

制作日期：2019年01月07日

# Kubernetes 安装部署操作手册

## 摘要

Kubernetes 单词起源于希腊语，是“舵手”或者“领航员”的意思，是“管理者”和“控制论”的根源。K8s 是把用 k 代替 8 个字符“ubernete”而成的缩写。首先，他是一个全新的基于容器技术的分布式架构领先方案。Kubernetes(k8s) 是 Google 开源的容器集群管理系统（谷歌内部:Borg）。在 Docker 技术的基础上，为容器化的应用提供部署运行、资源调度、服务发现和动态伸缩等一系列完整功能，提高了大规模容器集群管理的便捷性。

Kubernetes 是一个完备的分布式系统支撑平台，具有完备的集群管理能力，多扩多层次的安全防护和准入机制、多租户应用支撑能力、透明的服务注册和发现机制、内建智能负载均衡器、强大的故障发现和自我修复能力、服务滚动升级和在线扩容能力、可扩展的资源自动调度机制以及多粒度的资源配额管理能力。

**关键词：**Kubernetes；自我修复；滚动升级；动态伸缩

# 目录

总页数：19 页

1. 初识 kubernetes.....	1
1.1kubernetes 简介.....	1
1.2kubernetes 特征.....	1
1.3 核心组件.....	1
1.4 拓扑图.....	1
2. 集群.....	2
2.1 环境介绍.....	2
2.2 安装必要的组件.....	2
2.2.1 安装 docker.....	2
2.2.2 安装集群组件.....	3
2.2 使用 kubeadm 初始化集群.....	4
2.3 加入集群.....	4
3. 验证.....	5
3.1 验证集群节点.....	5
3.2 验证 pod 信息.....	5
4. 部署.....	5
4.1 deployment.....	5
4.2 Service.....	7
5. 升级.....	8
5.1 查看当前镜像版本.....	8
5.2 更新镜像.....	8
6. 回滚.....	9
6.1 查看历史版本.....	9
6.2 回滚版本.....	9
7. Dashborad.....	10
7.1 镜像准备.....	10
7.2 编写 dashborad.yaml.....	10
7.3 创建和修改.....	15
7.4 创建 ServiceAccount.....	16
7.5 登录 UI.....	17

## 1. 初识 kubernetes

### 1.1 kubernetes 简介

Kubernetes 是一个全新的基于容器技术的分布式架构领先方案。是 Google 内部集群管理系统 Borg 的一个开源版本。直到 2015 年 4 月，随着论文发布，才被众人熟知。Kubernetes 是一个开放的开发平台。不局限于任何一种语言，没有限定任何编程接口。是一个完备的分布式系统支撑平台。它构建在 docker 之上，提供应用部署、维护、扩展机制等功能，利用 Kubernetes 能方便地管理跨机器运行容器化的应用。

### 1.2 kubernetes 特征

1. 自主的管理容器，保证云平台中的容器按照用户的期望状态运行着。
2. 自动扩容，弹性伸缩。
3. 自动监控，删除出故障的应用。
4. 容器编排成组，并提供容器间的负载均衡。
5. 自我调度，自我管理。

### 1.3 核心组件

kubectl:客户端命令行工具，作为整个系统的操作入口。

kube-apiserver:以 REST API 服务形式提供接口，作为整个系统的控制入口。

kube-controller-manager:执行整个系统的后台任务，包括节点状态状况、Pod 个数、Pods 和 Service 的关联等。

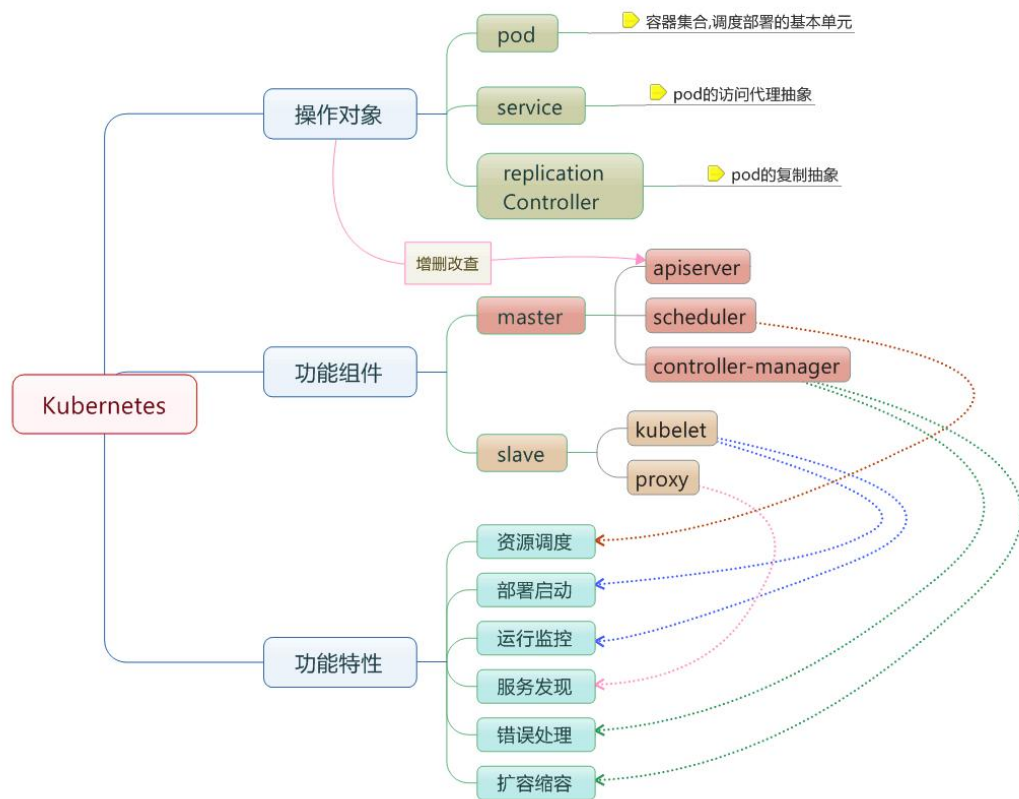
kube-scheduler:负责节点资源管理，接收来自 kube-apiserver 创建 Pods 任务，并分配到某个节点。

etcd:负责节点间的服务发现和配置共享。

kube-proxy:运行在每个计算节点上，负责 Pod 网络代理。定时从 etcd 获取到 service 信息来做相应的策略。

kubelet:运行在每个计算节点上，作为 agent，接收分配该节点的 Pods 任务及管理容器，周期性获取容器状态，反馈给 kube-apiserver。

### 1.4 拓扑图



## 2. 集群

### 2.1 环境介绍

192.168.2.5 k8s-master 4G 2 核

192.168.2.6 k8s-node 4G 2 核

### 2.2 安装必要的组件

#### 2.2.1 安装 docker

在每个节点执行

1. 添加阿里 yum 源

sudo yum-config-manager --add-repo

<http://mirrors.aliyun.com/docker-ce/linux/centos/docker-ce.repo>

2. 生成缓存

yum makecache fast

安装 docker-ce

```
yum -y install docker-ce
```

3.启动 docker 并设置为开机启动

```
systemctl start docker && systemctl enable docker
```

## 2.2.2 安装集群组件

每个节点执行

1. 修改 hosts 文件

vim /etc/hosts 添加以下信息

```
192.168.2.5 k8s-master
```

```
182.168.2.6 k8s-node
```

2.分别在不同主机上设置 hostname

```
hostnamectl --static set-hostname k8s-master
```

```
hostnamectl --static set-hostname k8s-node
```

3.关掉 selinux

```
setenforce 0
```

```
sed -i "s/^SELINUX=enforcing/SELINUX=disabled/g" /etc/sysconfig/selinux
```

4.关掉防火墙

```
systemctl stop firewalld
```

```
systemctl disable firewalld
```

5.关闭 swap

```
swapoff -a
```

```
sed -i 's/.*swap.*/#&/' /etc/fstab
```

6.配置转发参数

```
cat <<EOF > /etc/sysctl.d/k8s.conf net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1 EOF
```

```
sysctl --system
```

7.配置国内 yum 源并生成缓存

```
cat <<EOF > /etc/yum.repos.d/kubernetes.repo [kubernetes] name=Kubernetes
baseurl=https://mirrors.aliyun.com/kubernetes/yum/repos/kubernetes-el7-x86
_64/ enabled=1 gpgcheck=1 repo_gpgcheck=1
```

```
gpgkey=https://mirrors.aliyun.com/kubernetes/yum/doc/yum-key.gpg
```

```
https://mirrors.aliyun.com/kubernetes/yum/doc/rpm-package-key.gpg EOF
```

```
yum makecache fast
```

8.安装 kubeadm,kubectl,kubelet,ntpd

```
yum install -y kubeadm kubectl kubelet ntpdate
systemctl start kubelet && systemctl enable kubelet
```

## 2.2 使用 kubeadm 初始化集群

在 master 上面执行一下命令，初始化集群信息

```
kubeadm init --image-repository registry.aliyuncs.com/google_containers
--pod-network-cidr=10.244.0.0/16 --service-cidr=10.96.0.0/12
```

等待初始化完成后，系统提示我们执行一下三步：

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

并且初始化完成后，我们将 token 拷贝下来保存，后来加入 node 使用

```
kubeadm join 192.168.2.5:6443 --token 6pq12p.qolwnkesr8nu3ger
--discovery-token-ca-cert-hash
```

```
sha256:eb1dcad8d31fd03c4cf9ab0476c360da13b2204e01885ca8f98eafd43345ee8d
```

加入 flannel:

```
kubectl apply -f
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

**注意：**

这里有一个镜像拉取不下来我们手动拉取国内，并给改镜像打一个标签

```
docker pull registry.cn-hangzhou.aliyuncs.com/kubernetes_containers/flannel:v0.10.0-amd64
docker tag registry.cn-hangzhou.aliyuncs.com/kubernetes_containers/flannel:v0.10.0-amd64
quay.io/coreos/flannel:v0.10.0-amd64
```

## 2.3 加入集群

在 node 上面执行之前生成的 token

```
kubeadm join 192.168.2.5:6443 --token 6pq12p.qolwnkesr8nu3ger
--discovery-token-ca-cert-hash
sha256:eb1dcad8d31fd03c4cf9ab0476c360da13b2204e01885ca8f98eafd43345ee8d
```

## 3. 验证

### 3.1 验证集群节点

上面执行完命令后，集群不会马上变成 **ready** 状态，因为系统需要去下载 **docker** 镜像，稍等片刻后我们可以执行一下命令验证集群状态。

Kubectl get node

或者

Kubectl get node -o wide 显示更新信息

```
[root@k8s-master ~]# kubectl get node
NAME          STATUS    ROLES    AGE   VERSION
k8s-master    Ready     master   40m   v1.13.1
k8s-node      Ready     <none>    28m   v1.13.1
[root@k8s-master ~]# kubectl get node -o wide
NAME          STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
k8s-master    Ready     master   40m   v1.13.1   192.168.2.5   <none>         CentOS Linux 7 (Core) 3.10.0-693.el7.x86_64  docker://18.9.0
k8s-node      Ready     <none>    28m   v1.13.1   192.168.2.6   <none>         CentOS Linux 7 (Core) 3.10.0-693.el7.x86_64  docker://18.9.0
```

当所有节点都变成 **ready** 后，表示集群搭建完成。

### 3.2 验证 pod 信息

kubectl get pod -n kube-system

```
[root@k8s-master ~]# kubectl get pod -n kube-system
NAME                                READY   STATUS    RESTARTS   AGE
coredns-78d4cf999f-4r49q           1/1     Running   2           32m
coredns-78d4cf999f-q9h7s           1/1     Running   2           32m
etcd-k8s-master                     1/1     Running   2           32m
kube-apiserver-k8s-master           1/1     Running   2           32m
kube-controller-manager-k8s-master 1/1     Running   2           32m
kube-flannel-ds-amd64-rgpqg         1/1     Running   4           20m
kube-flannel-ds-amd64-vqnsf         1/1     Running   3           28m
kube-proxy-46dww                    1/1     Running   1           20m
kube-proxy-dcxfl                    1/1     Running   2           32m
kube-scheduler-k8s-master           1/1     Running   2           31m
```

当所有 pod 的 **ready** 都变成 **1/1** 就表示集群搭建成功。

## 4. 部署

### 4.1 deployment

一次 deployment 为一次部署，如图 Deployment 调用 ReplicaSet 创建多个 Pod 副本。而 ReplicaSet 不需要我们去管理，所以我们只需要创建一个 deployment 即可。我们编写一个 **nginx-deployment.yaml**。



```

apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
        track: stable
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80

```

apiVersion: 组名/版本号

Kind: 类型

Metadata: 元数据, 下面 name 表示当前 deployment 名称为 nginx-deployment

Spec: 规格

Replicas:3 表示副本集 3 个

Spec.template.metadata.labels 自定义标签, 一般配合 selector

spec.template.containers:可以有多个容器

- name 表示容器[1]的名字

Image: 使用镜像 nginx:1.7.9

containerPort: 容器端口 80

执行 `kubectl apply -f nginx-deployment.yaml` 即可创建 deployment

查看 deployment:

`kubectl get deploy/deployment/deployments`

查看 pod:

`Kubectl get pod`

此时 pod 显示正在创建, 因为镜像需要去 pull 下来。

```

[root@k8s-master kubernetes]# kubectl get pod
NAME                                READY   STATUS              RESTARTS   AGE
nginx-deployment-79694b8c57-bwxhs   0/1     ContainerCreating   0           97s
nginx-deployment-79694b8c57-cjg2k   0/1     ContainerCreating   0           97s
nginx-deployment-79694b8c57-ts6ht   0/1     ContainerCreating   0           97s
[root@k8s-master kubernetes]#

```

等待一段时间后, 我们的 pod 已经跑起来了。

```

[root@k8s-master kubernetes]# kubectl get pod
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-79694b8c57-bwxhs   1/1     Running   0           3m13s
nginx-deployment-79694b8c57-cjg2k   1/1     Running   0           3m13s
nginx-deployment-79694b8c57-ts6ht   1/1     Running   0           3m13s
[root@k8s-master kubernetes]#

```

此时我们使用 `-o wide` 查看 ip, 此时的应用只有集群内部能访问, 外部暂时不能访问。使用一个 ip 在任何一个节点都可以访问。

```
[root@k8s-master kubernetes]# kubectl get pod -o wide
NAME                                READY    STATUS    RESTARTS   AGE    IP            NODE    NOMINATED NODE    READINESS GATES
nginx-deployment-79694b8c57-bwxhs   1/1      Running   0           4m5s   10.244.1.4    k8s-node    <none>             <none>
nginx-deployment-79694b8c57-cjg2k   1/1      Running   0           4m5s   10.244.1.2    k8s-node    <none>             <none>
nginx-deployment-79694b8c57-ts6ht   1/1      Running   0           4m5s   10.244.1.3    k8s-node    <none>             <none>
[root@k8s-master kubernetes]# curl 10.244.1.4
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
body {
width: 35em;
margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif;
}
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

## 4.2 Service

Service 定义了一个服务的访问入口地址，前端的应用通过这个入口地址访问其背后的一组由 Pod 副本组成的集群实例，来自外部的访问请求被负载均衡到后端的各个容器应用上。Service 与其后端 Pod 副本集群之间则是通过 Label Selector 来实现对接的。

nginx-svc.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
  labels:
    app: nginx
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: 80
  selector:
    app: nginx
```


这里我们讲一下 selector,选择的是之前 deployment, labels 定义的标签。

创建 Service:kubectl apply -f nginx-svc.yaml

查看 service:kubectl get svc/service

```
[root@k8s-master kubernetes]# kubectl get svc
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes          ClusterIP   10.96.0.1     <none>         443/TCP          77m
nginx-service       NodePort    10.106.49.182 <none>         80:30565/TCP     7s
[root@k8s-master kubernetes]#
```

因为我们的 type 指定的是 NodePort 所以我们现在可以直接在浏览器使用任意节点 ip+30565 访问。并且我们指定的 pod 有 3 个，他每次会根据负载策略去访问三个 pod 中的任意一个。



## Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](http://nginx.org).  
Commercial support is available at [nginx.com](http://nginx.com).

*Thank you for using nginx.*

## 5. 升级

### 5.1 查看当前镜像版本

Kubectl get pod

Kubectl describe pod podName

看了看到 image: nginx:1.7.9

```
[root@k8s-master kubernetes]# kubectl get pod
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-79694b8c57-bwxhs   1/1     Running   0           17m
nginx-deployment-79694b8c57-cjg2k   1/1     Running   0           17m
nginx-deployment-79694b8c57-ts6ht   1/1     Running   0           17m
[root@k8s-master kubernetes]# kubectl describe pod nginx-deployment-79694b8c57-bwxhs
Name:                               nginx-deployment-79694b8c57-bwxhs
Namespace:                           default
Priority:                             0
PriorityClassName:                    <none>
Node:                                k8s-node/192.168.2.6
Start Time:                          Fri, 04 Jan 2019 16:52:36 +0800
Labels:                               app=nginx
                                      pod-template-hash=79694b8c57
                                      track=stable
Annotations:                          <none>
Status:                               Running
IP:                                   10.244.1.4
Controlled By:                        ReplicaSet/nginx-deployment-79694b8c57
Containers:
  nginx:
    Container ID:   docker://de24c5c58361619ca0fe609d2994a2135cf9c9aa698302f9cebc7ef450c59955
    Image:           nginx:1.7.9
    Image ID:        docker-pullable://nginx@sha256:e3456c851a152494c3e4ff5fcc26f240206abac0c9d794affb40e0714846c451
    Port:            80/TCP
    Host Port:       0/TCP
    State:           Running
      Started:       Fri, 04 Jan 2019 16:54:40 +0800
    Ready:           True
    Restart Count:   0
    Environment:     <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-9r5ws (ro)
Conditions:
```

### 5.2 更新镜像

Kubectl set image deploy/nginx-deployment nginx=nginx:1.10 --record

```
[root@k8s-master kubernetes]# kubectl set image deploy/nginx-deployment nginx=nginx:1.10 --record
deployment.extensions/nginx-deployment image updated
[root@k8s-master kubernetes]# kubectl get pod
NAME                                READY   STATUS             RESTARTS   AGE
nginx-deployment-79694b8c57-cjg2k   1/1     Running            0           22m
nginx-deployment-79694b8c57-ts6ht   1/1     Running            0           22m
nginx-deployment-86bc9bd78b-bpw6x   0/1     ContainerCreating  0           5s
nginx-deployment-86bc9bd78b-xqxp    0/1     ContainerCreating  0           5s
[root@k8s-master kubernetes]#
```

Kubectl get pod 可以看到之前的 pod 并不会被马上终止，而是一步步新运行一个，终止一个，达到灰度发布，不重启服务器更新的目的。稍等片刻后我们就可以看到之前的 3 个 pod 全部被替换。

```
[root@k8s-master kubernetes]# kubectl get pod
NAME                                READY    STATUS    RESTARTS   AGE
nginx-deployment-86bc9bd78b-9twrv   1/1      Running   0           109s
nginx-deployment-86bc9bd78b-bpw6x   1/1      Running   0           3m21s
nginx-deployment-86bc9bd78b-xqxp    1/1      Running   0           3m21s
[root@k8s-master kubernetes]# kubectl describe pod nginx-deployment-86bc9bd78b-9twrv
Name:                               nginx-deployment-86bc9bd78b-9twrv
Namespace:                          default
Priority:                             0
PriorityClassName:                   <none>
Node:                               k8s-node/192.168.2.6
Start Time:                         Fri, 04 Jan 2019 17:16:09 +0800
Labels:                             app=nginx
                                      pod-template-hash=86bc9bd78b
                                      track=stable
Annotations:                         <none>
Status:                             Running
IP:                                 10.244.1.7
Controlled By:                      ReplicaSet/nginx-deployment-86bc9bd78b
Containers:
  nginx:
    Container ID:                    docker://dcccffaba37090bc79a52d2ece92d9df6dbc0882b358a8cce1164c13787fb7061
    Image:                          nginx:1.10
    Image ID:                       docker-pullable://nginx@sha256:6202beb06ea61f44179e02ca965e8e13b961d12640101fca213efbfd145d75
    Port:                           80/TCP
    Host Port:                      0/TCP
    State:                          Running
      Started:                      Fri, 04 Jan 2019 17:16:10 +0800
    Ready:                          True
    Restart Count:                   0
    Environment:                    <none>
```

紧着我们查看 pod 信息，运行正常，版本也成功更新到 1.10

## 6. 回滚

### 6.1 查看历史版本

Kubectl rollout history deploy/nginx-deployment

```
[root@k8s-master kubernetes]# kubectl rollout history deploy/nginx-deployment
deployment.extensions/nginx-deployment
REVISION  CHANGE-CAUSE
1         <none>
2         kubectl set image deploy/nginx-deployment nginx=nginx:1.10 --record=true
[root@k8s-master kubernetes]#
```

可以看到以上有 1 和 2 两个版本，1 表示之前的，2 表示当前的，可能有 2,3,4 那么最新的就表示当前版本。接着我们回滚一下。

### 6.2 回滚版本

回滚到上一个版本：

Kubectl rollout undo deploy/nginx-deployment

回滚到指定版本：

Kubectl rollout undo deploy/nginx-deployment --to-revision=2

我们回滚到上一个版本试试，可以看到我们镜像版本已经回退到之前的 1.7.9 了。也是逐个替换，滚动更新。

```
[root@k8s-master kubernetes]# kubectl rollout undo deploy/nginx-deployment
deployment.extensions/nginx-deployment rolled back
[root@k8s-master kubernetes]# kubectl get pod
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-79694b8c57-5sk58   1/1     Running   0           10s
nginx-deployment-79694b8c57-pcng2   1/1     Running   0           10s
nginx-deployment-79694b8c57-srskc   1/1     Running   0           7s
nginx-deployment-86bc9bd78b-bpw6x   0/1     Terminating 0           8m49s
[root@k8s-master kubernetes]# kubectl get pod
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-79694b8c57-5sk58   1/1     Running   0           20s
nginx-deployment-79694b8c57-pcng2   1/1     Running   0           20s
nginx-deployment-79694b8c57-srskc   1/1     Running   0           17s
[root@k8s-master kubernetes]# kubectl describe pod nginx-deployment-79694b8c57-5sk58
Name:                               nginx-deployment-79694b8c57-5sk58
Namespace:                           default
Priority:                               0
PriorityClassName:                     <none>
Node:                                 k8s-node/192.168.2.6
Start Time:                           Fri, 04 Jan 2019 17:23:15 +0800
Labels:                               app=nginx
                                      pod-template-hash=79694b8c57
                                      track=stable
Annotations:                           <none>
Status:                               Running
IP:                                   10.244.1.8
Controlled By:                         ReplicaSet/nginx-deployment-79694b8c57
Containers:
  nginx:
    Container ID:   docker://e2b12bcdcb123d58501e71241ef7cd7402e07cf48ac8361c21db872d507cc89
    Image:          nginx:1.7.9
    Image ID:       docker-pullable://nginx@sha256:e3456c851a152494c3e4ff5fcc26f240206abac0c9d794affb40e0714846c451
    Port:           80/TCP
    Host Port:      0/TCP
```

## 7. Dashborad

### 7.1 镜像准备

因为 k8s 镜像都是需要从国外下载，我们需要先准备国内镜像。

注意：每个节点，包括 master 和 node 节点。

vim dashborad-image.sh

```
#!/bin/bash
```

```
docker pull registry.cn-shanghai.aliyuncs.com/qubit/kubernetes-dashboard-amd64:v1.10.0
```

```
docker          tag          registry.cn-shanghai.aliyuncs.com/qubit/kubernetes-dashboard-amd64:v1.10.0
```

```
k8s.gcr.io/kubernetes-dashboard-amd64:v1.10.1
```

```
docker image rm registry.cn-shanghai.aliyuncs.com/qubit/kubernetes-dashboard-amd64:v1.10.0
```

```
chmod +x image.sh ; sh image.sh
```

### 7.2 编写 dashborad.yaml

```
# Copyright 2017 The Kubernetes Authors.
```

```
#
```

```
# Licensed under the Apache License, Version 2.0 (the "License");
```

```
# you may not use this file except in compliance with the License.
```

```
# You may obtain a copy of the License at
```

```
#
```

```
# http://www.apache.org/licenses/LICENSE-2.0
```

```
#
```

# Unless required by applicable law or agreed to in writing, software  
# distributed under the License is distributed on an "AS IS" BASIS,  
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or  
implied.

# See the License for the specific language governing permissions and  
# limitations under the License.

# ----- Dashboard Secret ----- #

apiVersion: v1  
kind: Secret  
metadata:  
 labels:  
 k8s-app: kubernetes-dashboard  
 name: kubernetes-dashboard-certs  
 namespace: kube-system  
type: Opaque

---

# ----- Dashboard Service Account ----- #

apiVersion: v1  
kind: ServiceAccount  
metadata:  
 labels:  
 k8s-app: kubernetes-dashboard  
 name: kubernetes-dashboard  
 namespace: kube-system

---

# ----- Dashboard Role & Role Binding ----- #

kind: Role  
apiVersion: rbac.authorization.k8s.io/v1



```
metadata:
  name: kubernetes-dashboard-minimal
  namespace: kube-system
rules:
  # Allow Dashboard to create 'kubernetes-dashboard-key-holder' secret.
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["create"]
  # Allow Dashboard to create 'kubernetes-dashboard-settings' config map.
- apiGroups: [""]
  resources: ["configmaps"]
  verbs: ["create"]
  # Allow Dashboard to get, update and delete Dashboard exclusive secrets.
- apiGroups: [""]
  resources: ["secrets"]
  resourceNames: ["kubernetes-dashboard-key-holder",
"kubernetes-dashboard-certs"]
  verbs: ["get", "update", "delete"]
  # Allow Dashboard to get and update 'kubernetes-dashboard-settings' config
map.
- apiGroups: [""]
  resources: ["configmaps"]
  resourceNames: ["kubernetes-dashboard-settings"]
  verbs: ["get", "update"]
  # Allow Dashboard to get metrics from heapster.
- apiGroups: [""]
  resources: ["services"]
  resourceNames: ["heapster"]
  verbs: ["proxy"]
- apiGroups: [""]
  resources: ["services/proxy"]
  resourceNames: ["heapster", "http:heapster:", "https:heapster:"]
  verbs: ["get"]
```

---

apiVersion: rbac.authorization.k8s.io/v1

kind: RoleBinding

metadata:

name: kubernetes-dashboard-minimal

namespace: kube-system

roleRef:

apiGroup: rbac.authorization.k8s.io

kind: Role

name: kubernetes-dashboard-minimal

subjects:

- kind: ServiceAccount

name: kubernetes-dashboard

namespace: kube-system

---

# ----- Dashboard Deployment ----- #

kind: Deployment

apiVersion: apps/v1

metadata:

labels:

k8s-app: kubernetes-dashboard

name: kubernetes-dashboard

namespace: kube-system

spec:

replicas: 1

revisionHistoryLimit: 10

selector:

matchLabels:

k8s-app: kubernetes-dashboard

template:

metadata:

labels:



k8s-app: kubernetes-dashboard

spec:

containers:

- name: kubernetes-dashboard

image: k8s.gcr.io/kubernetes-dashboard-amd64:v1.10.1

ports:

- containerPort: 8443

protocol: TCP

args:

- --auto-generate-certificates

# Uncomment the following line to manually specify Kubernetes API

server Host

# If not specified, Dashboard will attempt to auto discover the API

server and connect

# to it. Uncomment only if the default does not work.

# - --apiserver-host=http://my-address:port

volumeMounts:

- name: kubernetes-dashboard-certs

mountPath: /certs

# Create on-disk volume to store exec logs

- mountPath: /tmp

name: tmp-volume

livenessProbe:

httpGet:

scheme: HTTPS

path: /

port: 8443

initialDelaySeconds: 30

timeoutSeconds: 30

volumes:

- name: kubernetes-dashboard-certs

secret:

secretName: kubernetes-dashboard-certs

- name: tmp-volume

```

        emptyDir: {}
        serviceAccountName: kubernetes-dashboard
        # Comment the following tolerations if Dashboard must not be deployed
on master
        tolerations:
        - key: node-role.kubernetes.io/master
          effect: NoSchedule

---
# ----- Dashboard Service ----- #

kind: Service
apiVersion: v1
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
  namespace: kube-system
spec:
  ports:
    - port: 443
      targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard

```

### 7.3 创建和修改

创建：Kubectl create -f dashborad.yaml

因为创建出来的 service 默认是 ClusterIP,只能集群内访问，修改一下类型需要改成 NodePort。

```
kubectl patch svc kubernetes-dashboard -p '{"spec":{"type":"NodePort"}}' -n kube-system
```

查看一下 service

```
Kubectl get svc -n kube-system
```

```
[root@k8s-master kubernetes]# kubectl get svc -n kube-system
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP	2d18h
kubernetes-dashboard	NodePort	10.104.86.29	<none>	443:31141/TCP	38m

```
[root@k8s-master kubernetes]#
```

这里我们对外暴露的端口是 31141，此时我们可以直接用 `https://master_ip:31141` 来访问，这个端口是随机的默认是 30000-32767

注意：我们只能用 firefox 浏览器

此时还不能直接访问，我们还需要添加一个授权管理的用户。

## 7.4 创建 ServiceAccount

```
vim admin-user.yaml
```

```
apiVersion: v1
```

```
kind: ServiceAccount
```

```
metadata:
```

```
  labels:
```

```
    k8s-app: kubernetes-dashboard
```

```
  name: admin
```

```
  namespace: kube-system
```

```
---
```

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: ClusterRoleBinding
```

```
metadata:
```

```
  name: admin
```

```
roleRef:
```

```
  apiGroup: rbac.authorization.k8s.io
```

```
  kind: ClusterRole
```

```
  name: cluster-admin
```

```
subjects:
```

```
- kind: ServiceAccount
```

```
  name: admin
```

```
  namespace: kube-system
```

创建： `kubectl create -f admin-user.yaml`

查看 svc: `kubectl describe serviceaccount admin -n kube-system`

此时会有一个 tokens 的描述

查看 `token:kubectl describe secret tokens 描述 -n kube-system`  
此时我们拷贝生成的 token。

## 7.5 登录 UI

此时输入 `https://master_ip:31141`

此时我们选择高级，添加例外，然后选择 token，将上面的 token 黏贴后即可登录。

